

STAGE latency comparison: gpt_pipeline

